

Project Phase Three Report

Yuning Hui – A53091905

Zhe Wang – A53097553

Test Environment

Apple Macbook Air with MacOS 10.11.5 and Postgres 9.5.

1 Cats

1.1 Option “Overall Likes”

The query plan shows that Postgres did not make use of any indices, including the automatic indices on primary keys.

Execution time: 170.877 ms

1.2 Option “Friend Likes”

The query plan shows that Postgres did not make use of any created indices. However, it uses the automatic generated indices on the primary keys of tables *like_activity* and *friendship*.

Execution time: 0.238 ms

1.3 Option “Friends-of-Friends Likes”

The query plan shows that Postgres did not make use of any created indices. However, it uses the automatic generated indices on the primary keys of tables *like_activity* and *friendship*.

Execution time: 1.602 ms

1.4 Option “My kind of cats”

The query plan shows that Postgres uses the created index on column *video_id* in table *like_activity*. It replaces a sequential scan by a bitmap heap scan on *like_activity*.

Execution time before creating index: 680.296 ms

Execution time after creating index: 645.572 ms

Performance improvement: 5%

1.5 Option “My kind of cats – with preference

The situation here is basically the same with the last query “My kinds of cats”. The query plan shows that Postgres uses the created index on column *video_id* in table *like_activity*. It replaces a sequential scan by a bitmap heap scan on *like_activity*.

Execution time before creating index: 757.017 ms

Execution time after creating index: 737.549 ms

Performance improvement: 2.7%

As we can see, in Cats, the only useful index is *like_activity(video_id)*. But it only gives a 3% ~ 6% performance. For other index choices, either they were not adopted by the query plan or they only bring minor improvements, which can be ignored. For example, an index on *like_activity(user_id)* only improves about 0.02% of the performance in the query *friend of friend likes*. As a conclusion, we recommend no index on Cats.

2 Sales

2.1 Total Sales

The query plan makes a great use of index *sales(customer_id)*. In the original query plan without the index, it materializes the table after sorting the entire table *sales*. The hash index can replace it with two index scans.

Execution time before creating index: 15175.472 ms

Execution time after creating index: 1308.359 ms

Performance improvement: 91.4%

2.2 Total Sales for Each State

The query plan makes no use of any indices created, including the automatic created indices on primary keys.

Execution time: 6866.951 ms

2.3 Total Sales for Each Product for a Given Customer

The query makes a great use of index *sales(customer_id)*. It avoids the sequential scan on table *sales*. Instead, it uses a bitmap index scan.

Execution time before creating index: 487.220 ms

Execution time after creating index: 0.136 ms

Performance improvement: 99.97%

2.4 Total Sales for Each Product and Customer Order by Dollar Value

The b-tree index *sales(product_id, customer_id)* can save an external merge sort. However, it only saves about 5% overhead.

Execution time before creating index: 29747.477 ms

Execution time after creating index: 28278.425 ms

Performance improvement: 4.9%

2.5 Total Sales for Each Product Category and State

The query plan shows that it uses no index at all.

Execution time: 11673.393 ms

2.6 For each one of the top 20 product categories and top 20 customers, it returns a tuple (top product, top customer, quantity sold, dollar value)

The index *sales(customer_id)* can save a sequential scan on table *sales*.

Execution time before creating index: 15782.326 ms

Execution time after creating index: 13961.283 ms

Performance improvement: 11.5%

We have tried many index combination choices, the only index which brings significant improvements is *sales(customer_id)*.